

COM3250 / COM6170

Introduction to Machine Learning

Instructor	Rob Gaizauskas Email: r.gaizauskas@dcs.shef.ac.uk Office: Room G28b @ Computer Science
Classes	Lecture 1: Monday 10:00 am, SG–LT08 (LT8, St. George's Complex) Tutorial/Lab: Monday 11:00 pm, SG—LT08 (LT8, St. George's Complex) Lecture 2: Tuesday 10:00 am, SG–LT08 (LT11, St. George's Complex)
Homepage	www.dcs.shef.ac.uk/~robertg/campus_only/com3250/
Assessment	Coursework: 30% Examination: 70 %
Reading Weeks	Weeks 6 and 12
Office Hours	Email requests for appointment

Slide 1

Course Aims and Objectives

Aims

- to describe the main approaches to automated concept learning
- to discuss the relationship between natural and artificial forms of learning
- to develop students' skills in designing and building serious artificial intelligence programs

Objectives

By the end of this course the students should:

- understand the main approaches that are used for representing concepts and learning them automatically;
- understand the relationships between natural learning processes and machine learning techniques;
- be able to develop effective software to implement common forms of knowledge representation;
- be able to develop effective software that will apply representative automated learning techniques using such representations.

Slide 2

Slide 3

Course Presuppositions

- knowledge of basic artificial intelligence techniques, from COM1080
- Advisable: some knowledge of machine reasoning techniques, from COM3290 (formerly COM2100).

Slide 4

Course Structure

- Introduction to Machine Learning (Lecture 1)
- Concept Learning (Lecture 2 and 3)
- Decision Tree Learning (Lecture 4 and 5)
- Evaluating Hypotheses (Lecture 6 and 7)
- Bayesian Learning, Bayesian Classifiers, Bayesian Belief Networks (Lecture 8 – 10)
- Instance-based Learning (Lecture 11 and 12)
- Rule Set Learning, Induction and Inductive Logic Programming (Lecture 13 – 16)
- Computational Learning Theory (Lecture 17 and 18)

Slide 5

References

Primary Textbook

T. Mitchell. *Machine Learning*. WCB/McGraw-Hill, Boston, 1997.

<http://www.cs.cmu.edu/~tom/mlbook.html>

Also Recommended

I. H. Witten and E. Frank. *Data Mining: Practical Machine Learning Tools and Techniques with Java Implementations*, Morgan Kaufmann, San Francisco, 2000.

Slide 6

Introduction to Machine Learning

Lecture Outline:

- What is Machine Learning?
- Why Study Machine Learning?
- Applications of Machine Learning
- Designing a Learning System: Overview
- WEKA: Software for Machine Learning

Reading:

Chapter 1 of Mitchell

Chapter 1 of Witten & Frank

Slide 7

What is Machine Learning ?

- A possible definition:

The study of how to design computer programs whose performance at some task improves through experience.

or more precisely (Mitchell):

Definition: A computer program is said to **learn**

- from experience E
 - with respect to some class of tasks T and
 - performance measure P
- if its performance at tasks in T as measured by P improves with experience E

- Important issue:

Are we only interested in *performance* of learning program?

Or are we also interested in discovering *human-comprehensible descriptions of patterns* in data? (*knowledge discovery*)

Slide 8

What is Machine Learning ? (cont)

- Disciplines Contributing to Machine Learning

Statistics

Biology

Artificial Intelligence

Cognitive Science

Philosophy

Computational Complexity

Information Theory

Control Theory

Why Study Machine Learning ?

- *technological or engineering motivation*
to build computer systems that can improve their performance at tasks with experience (data)
 - massive growth in on-line data
- *computer science motivation*
to understand better properties of various algorithms for function approximation
 - how the data must be represented
 - how much data they require
 - how accurate they can be
 - how to choose optimal data for training
- *cognitive science motivation*
to understand better how humans learn by modelling the learning process

Slide 9

Areas of Application of Machine Learning

- Data mining: using historical data to improve decisions – increasingly important given explosion of electronic data. E.g.:
 - *Medicine*: medical records → medical knowledge
 - * selecting best embryos from *in vitro* fertilisation based on 60 features of embryos and historical data on viability
 - *Business*: customer records → better business decisions
 - * assessing credit-worthiness of loan applicants based on features of former borrowers and repayment outcomes
 - * improving customer retention based on discovering patterns of features amongst loyal vs. defecting customers
 - *Agriculture* herd/crop records → better farming decisions
 - * improving cull selection from dairy herds by data mining over database of 700 attributes of millions of cows
 - ...

Slide 10

Slide 11

Areas of Application of Machine Learning (cont)

- Software applications we can't program by hand
 - autonomous driving
 - speech recognition
- Self customizing programs
 - Newsreader that learns user interests
- Computer Games
 - Heuristic evaluation functions for combinatorially explosive board games (chess, checkers, Othello)

Slide 12

Designing a Learning System: Overview

Task Design a checkers (draughts) program to play in world tournament.
(cf. Samuels, 1959; <http://www.cs.ualberta.ca/~chinook>)

Performance measure Percentage of games won in world tournament.

Principal design issues

- Choosing the training experience
- Choosing the target function
- Choosing a representation for the target function
- Choosing a function approximation algorithm

Designing a Learning System: Choosing the Training Experience

Choice of type of training experience has significant impact on success or failure of learning system

- Must decide between
 - *direct* training examples – e.g. individual checkers board states and correct moves for each; or
 - *indirect* training examples – e.g. move sequences and final outcomes of games.Here learner must additionally decide which moves in the sequence are good/bad (the *credit assignment*) problem.

Direct training examples easier to learn from but harder/more expensive to obtain

- Must decide how much learner controls training examples
 - teacher suggests board states and correct moves
 - learner asks teacher about novel/confusing board states
 - learner plays itself (no teacher)
 - random process outside learner provides examples
 - learner autonomously explores environment to collect training examples

Slide 13

Designing a Learning System: Choosing the Training Experience (cont)

- Must decide how well training examples represent distribution of examples over test space
 - ideally want training/test distributions to be identical, but not always possible in practice

- Returning to checkers problem:

Task T : playing checkers

Performance measure P : percentage of games won in world tournament

Training experience E : games played against itself (i.e. *indirect* training examples)

Slide 14

Designing a Learning System: Choosing the Target Function

- Must decide *what* will be learned and *how* this will be used by the *performance program*
- For the checkers program, assume program can generate legal moves from any given board state.
Want the program to learn the *best* move.
- Natural to think of this as learning a function

$$\text{ChooseMove} : B \rightarrow M$$

from set of legal board positions B to set of legal moves M

In machine learning such a function is called a *target function*.

Choosing the function to learn is a key design choice.

- Given the indirect training examples available, *ChooseMove* is very difficult to learn.
Better choice is a function

$$V : B \rightarrow \mathcal{R}$$

which assigns a *numerical value* to each board state – higher values = better board states

- Our system generates the successor board state for each legal move from the current state, then uses V to evaluate these states and picks the best

Slide 15

Designing a Learning System: Choosing the Target Function (cont)

- Continuing checkers example, choose a specific target function V

For any board state b :

1. $V(b) = 100$, if b is a final board state that is won
2. $V(b) = -100$, if b is a final board state that is lost
3. $V(b) = 0$, if b is a final board state that is drawn
4. $V(b) = V(b')$, if b is not a final board state, where b' is the best final board state that can be achieved by playing optimally from b , assuming opponent plays optimally too.

- This recursive definition specifies V for every $b \in B$.

However, V so defined is not efficiently computable, since for case 4 it requires searching every possible line of play till the end of the game.

Such a definition is called *non-operational*; we need an *operational* definition – one that can be used to evaluate states/select moves in realistic times.

- Perfectly learning an operational form of V usually very difficult – in general only possible to learn an *approximation* of V .

Call the function our system actually learns \hat{V} to distinguish it from the ideal target function V

Slide 16

Designing a Learning System: Choosing a Representation for the Target Function

- Many possible representations for target function \hat{V}
 - lookup table for every board state
 - rules that match against features of board states
 - neural network trained on features of board states
 - ...
- Tradeoff between
 - expressivity of representation – allows closer approximation to V ; and
 - volume of training data required – more expressive representations require more training data
- As a simple example, identify features:

x_1 : number of black pieces on board	x_2 : number of red pieces on board
x_3 : number of black kings on board	x_4 : number of red kings on board
x_5 : number of black pieces threatened by red	x_6 : number of red pieces threatened by black

and let \hat{V} be defined as a linear function of these features:

$$\hat{V}(b) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6$$

where $w_0 \dots w_6$ are weights to be learned by the learning algorithm.

Designing a Learning System: Choosing a Representation for the Target Function (cont)

Design of checkers learning program so far can be summarised as:

Task T : playing checkers

Performance measure P : percentage of games won in world tournament

Training experience E : games played against itself

Target function : $V : \text{Board} \rightarrow \mathcal{R}$

Target function representation :

$$\hat{V}(b) = w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + w_5x_5 + w_6x_6$$

**Designing a Learning System:
Choosing a function approximation algorithm**

- To learn \hat{V} need a set of training examples of the form

$$\langle b, V_{train}(b) \rangle$$

where b is a description of a board state in terms of the features $x_1 \dots x_6$, and $V_{train}(b)$ is the training value we wish to associate with that board state.

E.g.

$$\langle \langle x_1 = 3, x_2 = 0, x_3 = 1, x_4 = 4, x_5 = 0, x_6 = 0 \rangle, +100 \rangle$$

- To estimate training values for intermediate board states, use current estimate of \hat{V} and set

$$V_{train}(b) \leftarrow \hat{V}(Successor(b))$$

Note that \hat{V} is being used to estimate training values that will be used to refine \hat{V}

- the current estimate for $Successor(b)$ is being used to estimate the values for b
- OK, since estimates will be quite accurate near end of game and accuracy will be iteratively propagated back to earlier board states by the algorithm.

Slide 19

**Designing a Learning System:
Choosing a function approximation algorithm**

- Use the training examples to update the weights in \hat{V}
 - initialise weights to random values or set all equal
- Common approach is to modify weights by minimising the squared error E between the training values and those predicted by the hypothesis \hat{V} .

$$E \equiv \sum_{\langle b, V_{train}(b) \rangle \in \text{training examples}} (V_{train}(b) - \hat{V}(b))^2$$

- One algorithm for this is Least Mean Square (LMS) algorithm.

LMS weight update rule

For each training example $\langle b, V_{train}(b) \rangle$

- Use current weights to calculate $\hat{V}(b)$
- For each weight w_i , update as

$$w_i \leftarrow w_i + \eta(V_{train}(b) - \hat{V}(b))x_i$$

where η is a small constant (e.g. 0.1) that controls size of weight update

Slide 20

Designing a Learning System: Choosing a function approximation algorithm (cont)

- By choosing the weight update formula

$$w_i \leftarrow w_i + \eta(V_{train}(b) - \hat{V}(b))x_i$$

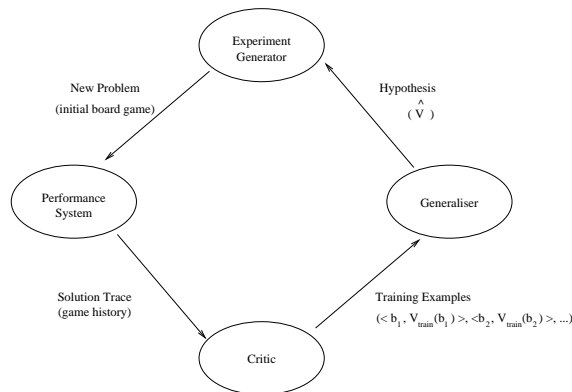
note:

- when $(V_{train}(b) - \hat{V}(b)) = 0$ no weight is changed
- when $(V_{train}(b) - \hat{V}(b)) > 0$, i.e. $\hat{V}(b)$ is too low, then w_i is increased proportionally to x_i – increases $\hat{V}(b)$, decreases E
- when $x_i = 0$ no change occurs to w_i , so weights are only updated for features which occur in training examples

In certain settings LMS can be proved to converge to least squared error approximation to V_{train} values.

Slide 21

Overall Design of Learning System



Performance System Performs the task (e.g. checkers) using learned target function (\hat{V}). I.e takes a new problem instance (new game) and proposes a solution (game history).

Critic Generates new training examples from solution trace, e.g. using $V_{train}(b) \leftarrow \hat{V}(Successor(b))$ estimates.

Generaliser Hypothesises target function (\hat{V}) from training examples, e.g. using LMS algorithm.

Experiment Generator Generates a new problem instance for the performance system, given the current hypothesis.

Slide 22

Slide 23

Summary

- Machine learning is about creating programs which improve their performance on a given task with experience
- Aside from performance alone, we may also be interested in discovering patterns in our data which we can understand
- ML is increasingly important in applications in all walks of life given the exponential increase in electronic data
- ML also of interest to theoretical computer scientists and to cognitive scientists interested in modelling human learning

Slide 24

Summary (cont)

- The general ML setting involves:
 - identifying the task, performance measure and experience (e.g. checkers, % games won, past games)
 - determining whether the training experience will be direct/indirect, the role of the teacher in training, the representativeness of the training data (e.g. games played against self)
 - choosing the target function to be learned, and if it is not practicably learnable, choosing an approximation of the ideal target function instead (e.g. $V : B \rightarrow \mathcal{R}$ instead of $ChooseMove : B \rightarrow M$)
 - choosing a representation of the target function which looks promising (e.g. some set of attributes describing an instance of interest – # red pieces, black pieces, kings, etc.)
 - choosing an algorithm which approximates the target function (e.g. LMS)



WEKA: Software for Machine Learning

- For labs and assignments we will be using **Weka** – a software platform for experimenting with and learning about machine learning algorithms (complements Witten & Frank book)
- Weka 3 is available from the University of Waikato in New Zealand at:
<http://www.cs.waikato.ac.nz/ml/weka>
- If you have your own laptop you are advised to download and install a copy.
- Weka-3.4.4 is already available on the DCS Linux Desktop, under the directory
`/usr/local/pkg/weka-3.4.4`
- Have a play! (see sample data files in data directory that comes with it)

Slide 25